AWS Disaster Recovery

Preventing unforeseeable disasters that cause business outages and product loss is one of the most significant challenges for IT experts today. The AWS Disaster Recovery Plan (DRP) helps you store and restore data to minimize disasters that may cause loss of infrastructure, plans, and data.

While there are many DR strategy initiatives in the market, AWS offers several service options within its own ecosystem that ensure business continuity.

Key factors for Disaster Planning

1. **Recovery Time Objective (RTO):** This marks the time it takes for the business to restore to its defined service level after a disaster has occurred. The business decides on the maximum acceptable time gap between interruption and restoration of the services.
2. **Recovery Point Objective (RPO):** This marks the acceptable amount of data loss in terms of time. The business decides the maximum acceptable amount of time since the last data recovery point.

AWS Disaster Recovery Methods

- **Backup and Restore**: This method is a simple and straightforward way to backup and restore data as needed, however, it can be time-consuming due to none of your data being on standby.

- **Pilot Light**: This method keeps critical applications and data at-the-ready so that it can be quickly retrieved if needed.

- **Warm Standby**: This method always keeps a duplicate version of your business' core elements running on standby, which makes for a little downtime and an almost seamless transition.
- **Multi-Site Solution**: This method fully replicates the data/applications between two or more active locations and splits your traffic between them. Interruptions simply trigger traffic to be rerouted to the unaffected area, resulting in almost zero downtime. But running two separate environments incurs much higher costs.